



BITS Pilani
Pilani Campus

Network Programming

K Hari Babu
Department of Computer Science & Information Systems



Course Overview

Instructor



- Dr K Hari Babu, Dept. of CS & IS
- khari@pilani.bits-pilani.ac.in
- 01596-51-5813

Objectives



- To study the Unix/Linux API related to
 - Process management
 - File I/O
 - Memory Management
 - Network
- To develop network-aware applications.
 - Unicast/Multicast/Broadcast application development
 - Client design & implementation: Browsers etc.
 - Server design & Implementation: web servers etc
 - Server scalability

Course Topics



- System Programming
 - Process, thread and signals management
 - File System API
 - User, User Groups, Process credentials
 - Daemons
 - Inter Process Communication
- Socket Programming
 - Sockets API
 - Writing client-server applications
 - Broadcasting and multicasting
 - Raw sockets and datalink access

Course Topics



- Client & Server Design Alternatives
 - I/O Models
 - Blocking, Non-blocking, I/O Multiplexing, Signal driven, Asynchronous I/O
 - Client design Models
 - Server design Models
- Distributed Programming
 - SUN RPC
 - XDR
 - DCE RPC
 - XML RPC
 - CORBA
 - DCOM

- Prescribed Text Books

- **T1:** Stevens, R.W., “Unix Network Programming, Vol-I Networking APIS : Sockets and XTI”, Prentice- Hall of India, 3rd Edition, 2004
- **T2:** Stevens, R.W., “Unix Network Programming: Vol-II Inter Process Communications”, Prentice-Hall of India, 2nd Edition, 1999.

- Reference Books

- **R1:** The Linux Programming Interface: Linux and UNIX System Programming Handbook by Michael Kerrisk, No Starch Press © 2010
<http://library.books24x7.com/toc.aspx?bookid=41558>
- **R2:** W.R. Stevens, Advanced Programming in the UNIX Environment, Addison Wesley, 1998.

Pre-requisites



- Course will mainly focus on Linux 2.6 kernel
- Programs will be in C language.
 - We will discuss Java API for sockets which provides a simpler and higher level abstraction
 - We will also discuss Windows API for sockets. But major focus will be on Linux.
- You are expected to be familiar with C.
- Some familiarity with Networks and OS concepts will help.
 - Not a absolute requirement.

Course Delivery



- Weekly 1 Lecture
 - Expecting participation & interactive sessions
- Labs
 - Total 8 labs on various topics
 - Helps you to make your hands dirty
 - <http://lab.bits-pilani.ac.in/courses/>
- Discussion Group
 - <http://taxila.bits-pilani.ac.in/>
 - Can post your queries
 - Will be checking frequently (twice a week) and replying you.

Evaluation Components



- Programming Assignments (15%)
 - No groups. Individual.
 - Assignment1
 - System Programming
 - Socket Programming
 - Assignment2
 - Client & Server Design
 - Distributed Programming
 - Malpractice
 - All parties will be equally penalized. Zero marks to all parties involved.
 - No sharing of source code.
- Mid Term (35%)
- Final (50%)



History of UNIX & C

A Brief History of Time (UNIX and C)



- 1969 – First Unix Ken Thompson at AT&T Bell Labs
 - Ideas from Multics:
 - Tree structured file system
 - Program for interpreting commands (shell)
 - Files – unstructured streams of bytes
- 1970 Unix rewritten in assembly for DEC PDP-11
- C – Dennis Ritchie – a systems programming language
 - BCPL → B (Thompson) → C
- 1973 Kernel rewritten in C – eases porting to other machines.
- 1977 Bell Labs released Unix System III
- 1982 AT & T released Unix System V.

Berkeley Software Division (BSD)



- (1975) Thompson visiting Prof. at UC-Berkeley
- A student Bill Joy added new features
 - Vi editor
 - C shell
 - First paging virtual memory management (Unix) BSD 4.2
 - Sendmail, Pascal compiler
 - Later co-founded Sun Microsystems
- Berkely Software Distribution (BSD)
 - Released BSD 3 in 1979
 - Latest one in 1994 BSD 4.4
- Derivatives
 - Darwin, Free BSD, Net BSD, Open BSD

Unix after 1979



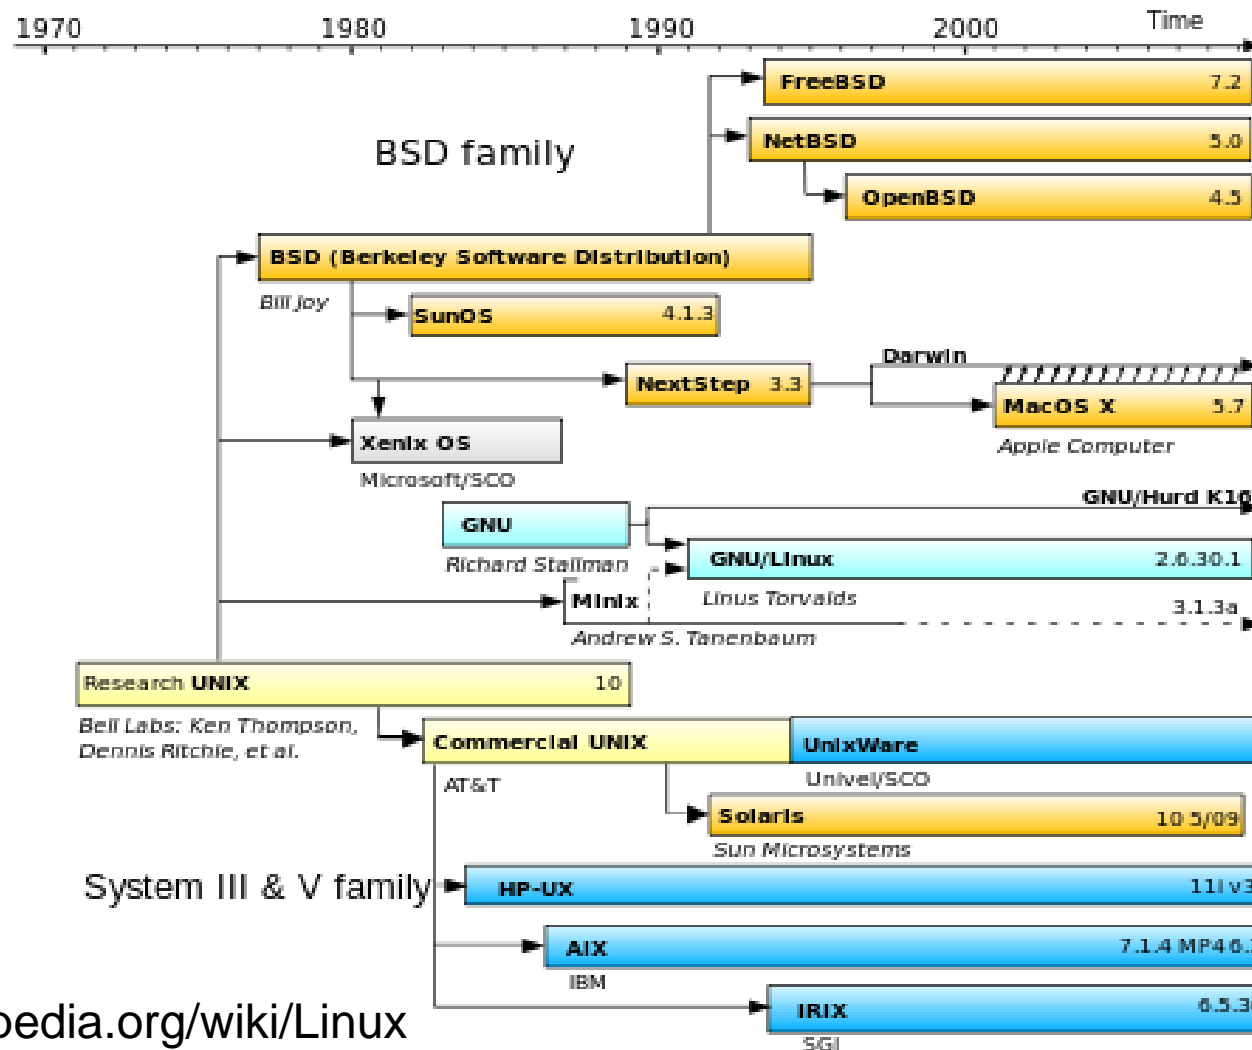
- BSD continued at UC-Berkeley
- Bell Labs System III → Systems V
- POSIX standard (1988)
- Other Software
 - X windows
 - Free Software Foundation
 - GNU Public License
- Minix – (1988)
 - Unix like; MINI-uniX; for education; A. Tannenbaum

Unix Commercial Versions



- In mid 90s many companies supported high-end features and ported them to their own architecture.
 - Digital's Tru64
 - HP-UX
 - IBM AIX
 - Sequent's DYNIX
 - SGI's IRIX
 - Sun Solaris
 - SunOS

Unix Timeline



<http://en.wikipedia.org/wiki/Linux>

Unix Characteristics/Strengths



- Unix is Simple
 - Has only hundreds of systems calls where as other OSs have thousands.
 - Has clear design goals
- Everything is a file
 - Same system calls `open()`, `read()`, `write()` and `close()` can work with files, devices, networks sockets.
- Unix kernel and system utilities are written in C.
 - Easy portability to different architectures.
- Fast process creation time.
- Robust inter process communication (IPC).

- Richard Stallman (1983) Goal a free Unix
 - Known for Free Software movement, GNU, Emacs, gcc
 - Never really released GNU operating system
- Free Software Foundation
 - <http://www.fsf.org/>



<http://en.wikipedia.org/wiki/GNU>

- Minix – (1988)
 - Unix like; MINI-uniX; for education; A. Tannenbaum
- (1991) Linus Torvalds
- For Intel x86 systems
- Moved to big Iron
- more than 90% of today's 500 fastest supercomputers run some variant of Linux
- Network routers
- Embedded systems
- Android


<http://en.wikipedia.org/wiki/Linux>



Overview of Linux OS Concepts

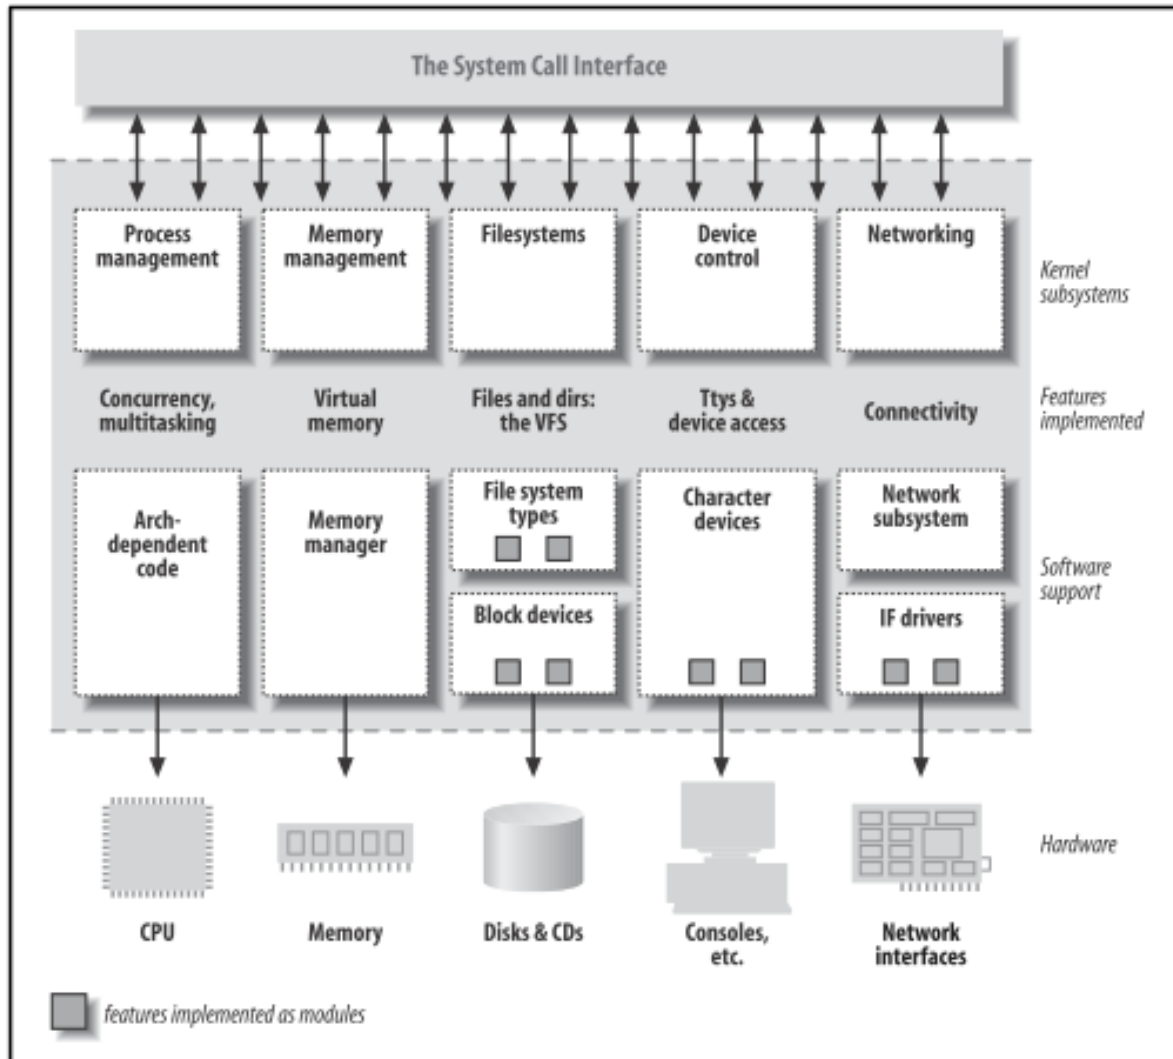
- Kernel
- C library
- Toolchain
- Basic system utilities
- Shell

About kernel



Source: Wikipedia

Kernel



Source: Linux Device Drivers

- Central software that manages and allocates computer resources (CPU, RAM, devices etc).
- Kernel greatly simplifies writing programs and using other programs.
 - Processes
 - Virtual address space
 - Virtual CPU

Kernel Tasks




- Process scheduling
- Memory management
- Provision of a file system
- Creation and termination of processes
- Access to devices
- Networking
- Provision of a system call application programming interface (API)

Memory Management



- Software size is enormous and RAM is a limited resource.
- Virtual Memory Management
 - Processes are isolated from one another and from the kernel.
 - Only a part of the process needs to be kept in memory. So multiple processes can be held in main memory simultaneously.
- Virtual Address Space
 - On 32 bit systems, there can be 4 GB address space.
 - Normally 2 GB is meant for Kernel and 2 GB for user process.
 - Current process has all the user address space with it.
 - Programmer flexibility.

Memory Management



Source: wikipedia

- Process Scheduling
 - Preemptive Multitasking
 - Multiple processes can simultaneously reside in memory
 - Preemptive means kernel decides which process and how long it gets CPU not the process itself.
- File System
 - File creation, retrieval, deletion etc.
 - More later
- Networking
 - TCP/IP stack
 - Sending/receiving messages on behalf of user processes.
 - More later

Kernel Tasks: Process Management




- Kernel can load new program into memory, provide it with the resources.
 - `fork()`
 - `execve()`
- Instance of a running program is termed as a process.
- Processes need to invoke system calls to access resources.
- Kernel executes the system calls on behalf of the process in kernel space.
- **Access to devices:**
 - Kernel provides an interface that standardizes and simplifies access to devices.
 - Keyboard, display etc.

Kernel Tasks: System Calls



- Kernel entry points are known as system calls.
- System calls are software interrupts.




Kernel Mode vs User Mode



- Processor architecture allows the CPU to operate in at least two different modes:
 - User mode
 - Kernel mode
- Similarly virtual memory is marked as user space and kernel space.
 - When running user mode, the CPU can access only user space.
 - When running in kernel mode, the CPU can access both user and kernel space.
- By this arrangement, user processes can't access instructions and data structures of kernel.

Kernel Space vs User Space

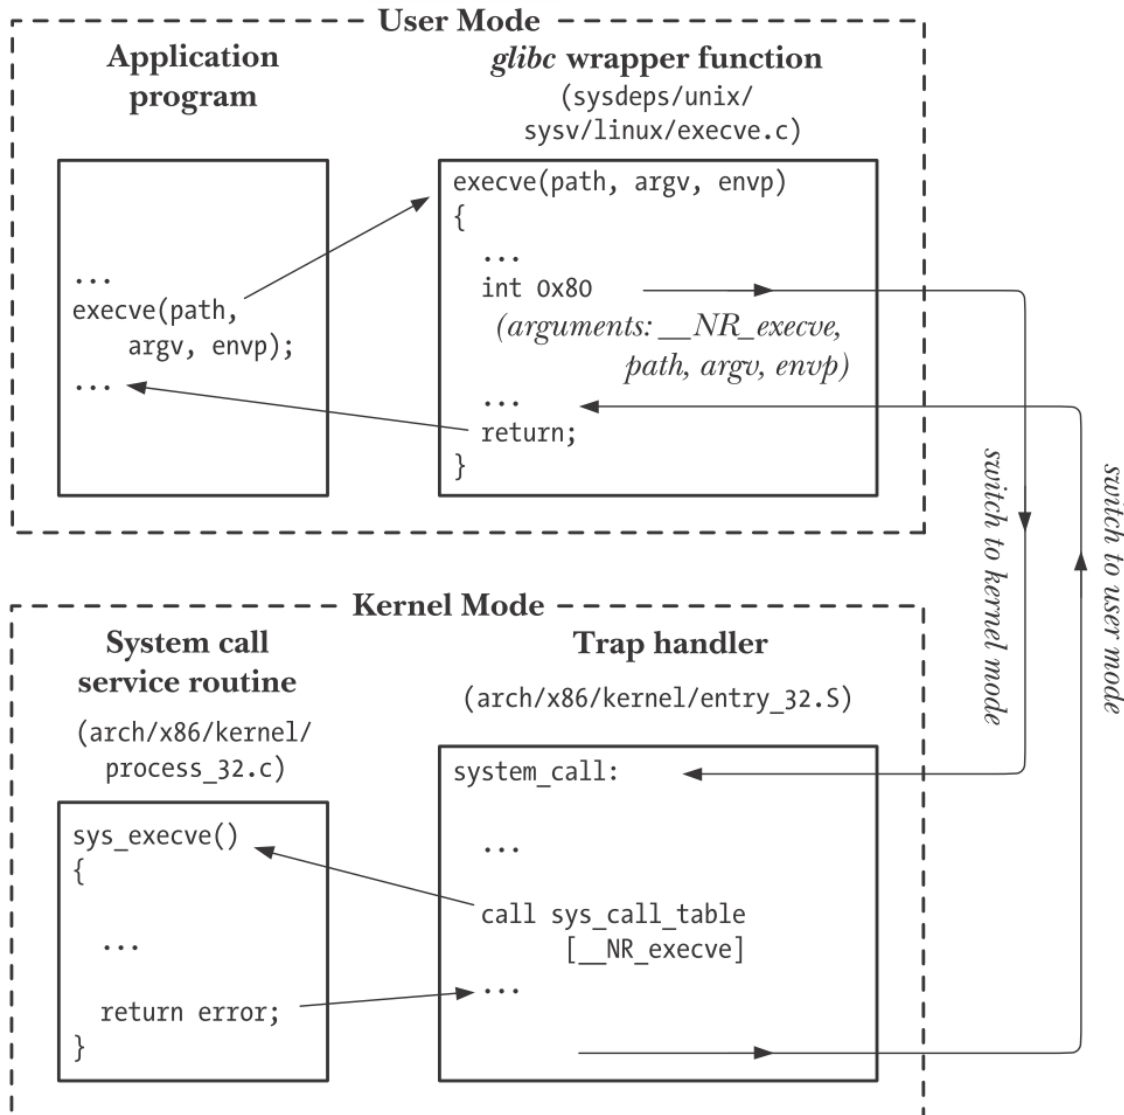


Kernel Tasks: System Calls



- A system call changes the processor from user mode to kernel mode.
- The set of sys calls are fixed. Each system call is identified by a unique number.
- System calls arguments are copied from user space to kernel space (registers).
- Wrapper function executes a trap instruction (int 0x80) causing the CPU to enter kernel mode.
- Example: `execve()`
 - System call number is 11.
 - `Sys_call_table` contains the address of the routine `sys_execve()`.
- Not all C library calls invoke system calls. E.g.: `<string.h>`

Kernel Tasks: System Calls



Q&A



Next Time



- Please read through R1: chapters 4-7



BITS Pilani
Pilani Campus



Thank You