



BITS Pilani
Pilani Campus

Network Programming

K Hari Babu
Department of Computer Science & Information Systems

Outline



- C Library
- Error Handling
- File I/O



BITS Pilani
Pilani Campus



Sys Calls & C Library

Sys Calls & C Library



- System call is a controlled entry point into the kernel.
- For every sys call there is a wrapper function in C library.
 - All library functions are not sys calls.
 - We use wrapper functions in programs
 - `fork()`, `execve()` ...
- Wrapper function copies the arguments into specific registers. Also copies sys call number into a specific CPU register.
- Wrapper functions executes trap instruction `int 0x80`.
 - Causes CPU to switch from user mode to kernel mode.
- Kernel executes *system_call* routine at *arch/x86/kernel/entry.S*

Sys Calls & C Library



- On i386, the parameters of a system call are transported via registers.
 - The system call number goes into %eax
 - the first parameter in %ebx
 - the second in %ecx
 - the third in %edx
 - the fourth in %esi
 - the fifth in %edi
 - the sixth in %ebp.

Sys Calls & C Library



- *system_call()* routine:
 - Saves register values onto the stack
 - Checks the validity of the system call number.
 - Invokes the corresponding service routine.
 - Service routine returns a result status to the *system_call* routine.
 - Restores register values from the kernel stack and places the system call return value on the user process stack.
 - Returns to the wrapper function, simultaneously returning the process to enter user mode.
- Wrapper function checks the return value and if it is an error it sets the *errno* variable.
- System calls incur an appreciable overhead.
 - Calling a C library wrapper function is synonymous with invoking the corresponding system call routine.

- Many library functions do not make use of system calls.
- Often library functions provide a more caller-friendly interface than the underlying sys call.
 - `fopen()` uses `open()`
 - `printf` uses `write()`
 - `malloc()` uses `brk()`
- GNU C library (*glibc*)
 - `libc.so.6`
 - Where is it stored? List dynamic dependencies
 - `ldd a.out`
 - Finding the version
 - `./libc.so.6`



BITS Pilani
Pilani Campus



Handling Errors

Errors from System Calls



- A service routine in kernel returns a negative number in case of error. The negative number corresponds to standard error codes.

`/usr/include/asm-generic/errno-base.h`

```
#ifndef _ASM_GENERIC_ERRNO_BASE_H
#define _ASM_GENERIC_ERRNO_BASE_H

#define EPERM          1  /* Operation not permitted */
#define ENOENT          2  /* No such file or directory */
#define ESRCH          3  /* No such process */
#define EINTR          4  /* Interrupted system call */
#define EIO             5  /* I/O error */
#define ENXIO          6  /* No such device or address */
#define E2BIG          7  /* Argument list too long */
#define ENOEXEC        8  /* Exec format error */
#define EBADF          9  /* Bad file number */
#define ECHILD        10  /* No child processes */
#define EAGAIN        11  /* Try again */
#define ENOMEM        12  /* Out of memory */
#define EACCES        13  /* Permission denied */
#define EFAULT        14  /* Bad address */
#define ENOTBLK       15  /* Block device required */
```

Errors from System Calls



- In case of error, wrapper function stores the positive value of the error code into `errno` variable and returns -1.

```
2 cnt = read(fd, buf, numbytes);
3 if (cnt == -1) {
4     if (errno == EINTR)
5         fprintf(stderr, "read was interrupted by a signal\n");
6     else {
7         /* Some other error occurred */
8     }
9 }
10
```

- `perror()` and `strerror()` can be used to print the error.

```
2 fd = open(pathname, flags, mode);
3 if (fd == -1) {
4     perror("open");
5     exit(EXIT_FAILURE);
6 }
7
```

errno variable



- It is present one per each process.
- It is set in wrapper function after sys call returns error.
 - Every time it is over written.
 - So only after a sys call returns -1, we refer to errno.

Handling Errors from Library Functions



- Some library functions return error information exactly like system calls.
 - Return -1
 - errno is set
- Some return value other than -1 but set errno
 - fopen
- Others do not use errno at all
 - gethostbyname

Tracing System Calls



- *strace* command allows us to trace the system calls made by a program.

```
haribabuk@haribabuk-VirtualBox ~ $ strace ./a.out
execve("./a.out", ["/a.out"], [/* 48 vars */]) = 0
brk(0)                                = 0x10e9000
access("/etc/ld.so.nohwcap", F_OK)     = -1 ENOENT (No such file or directory)
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fd26b5e2000
access("/etc/ld.so.preload", R_OK)     = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY)     = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=128950, ...}) = 0
mmap(NULL, 128950, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7fd26b5c2000
close(3)                               = 0
```

- Each system call is displayed with input and output arguments
- Arguments are printed in symbolic form.
- *ltrace* is used for tracing library calls.

```
haribabuk@haribabuk-VirtualBox ~ $ ltrace ./a.out
__libc_start_main(0x4005f4, 1, 0x7fffd10e0ba8, 0x400700, 0x400790 <unfinished ...>
socket(2, 3, 1)                                = -1
recvfrom(0xffffffff, 0x7fffd10e04d0, 1500, 0, 0) = -1
+++ exited (status 0) +++
```



File I/O

File Descriptor (fd)



- File Descriptors refer to all types of open files.
 - Pipes, FIFOs, sockets, terminals, devices, and regular files.
- Each process has its own set of file descriptors.
- All system calls refer to file descriptors for performing I/O.
- Three standard file descriptors.

File descriptor	Purpose	POSIX name	<i>stdio</i> stream
0	standard input	STDIN_FILENO	<i>stdin</i>
1	standard output	STDOUT_FILENO	<i>stdout</i>
2	standard error	STDERR_FILENO	<i>stderr</i>

- These three descriptors are open in the shell process.
 - Whenever a new program is executed in the shell, a child process is created. All three descriptors remain open in the child.
- File descriptors are different from FILE streams. FILE stream is a C library abstraction over fd.

File I/O: open()



```
2 #include <sys/stat.h>
3 #include <fcntl.h>
4 int open(const char * pathname , int flags , ... /* mode_t mode */);
5 Returns file descriptor on success, or -1 on error
```

- *flags* is a bitmask that refers to read only or write only or both.
- *mode* refers to permissions. *mode* is used only when creating a file.

```
2 openFlags = O_CREAT | O_WRONLY | O_TRUNC;
3 filePerms = S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP |
4             S_IROTH | S_IWOTH;          /* rw-rw-rw- */
5 outputFd = open(argv[2], openFlags, filePerms);
6 if (outputFd == -1)
7     errExit("opening file %s", argv[2]);
```

- O_CREAT option is used when a new file is to be created.
- O_TRUNC option is used when the data in the file has to be deleted.
- O_APPEND is used for appending to the existing file.
- Several other flags also present ... (R1: 4.3.1)

File I/O: read()



```
2  #include <unistd.h>
3  ssize_t read(int fd , void * buffer , size_t count );
4  /*Returns number of bytes read, 0 on EOF, or -1 on error*/
```

- Reads at most *count* bytes from the open file referred to by *fd* and stores them in a *buffer*.
- It returns the no of bytes actually read or EOF or -1.
 - *count*: maximum number of bytes to read
 - *buffer*: address of the memory buffer into which the input data is to be placed.
 - Read may read less than *count*.
 - In regular files, we may be close to EOF.
 - In pipes, FIFOs, and sockets it may read less than *count* due to non-availability of data.

File I/O: read()



```
2 char buffer[MAX_READ + 1];
3 ssize_t numRead;
4 numRead = read(STDIN_FILENO, buffer, MAX_READ);
5 if (numRead == -1)
6     errExit("read");
7 buffer[numRead] = '\0';
8 printf("The input data was: %s\n", buffer);
```

- STDIN_FILENO refers to fd 0.
- At line 9, we need to include NULL character because *read()* doesn't do it itself.

File I/O: write()



```
2  #include <unistd.h>
3  ssize_t write(int fd, void * buffer , size_t  count );
4  /*Returns number of bytes written, or -1 on error */
```

- Writes up to *count* bytes from *buffer* to the open file referred to by *fd*.
- Returns the number of bytes actually written which may be less than *count*.

File I/O: close()



```
2  #include <unistd.h>
3  int close(int  fd );
4  /*Returns 0 on success, or -1 on error*/
```

- It is called after all I/O has been completed in order to release the file descriptor *fd* and its associated kernel resources.
- When a process terminates all of its open file descriptors are automatically closed.

File I/O: lseek()



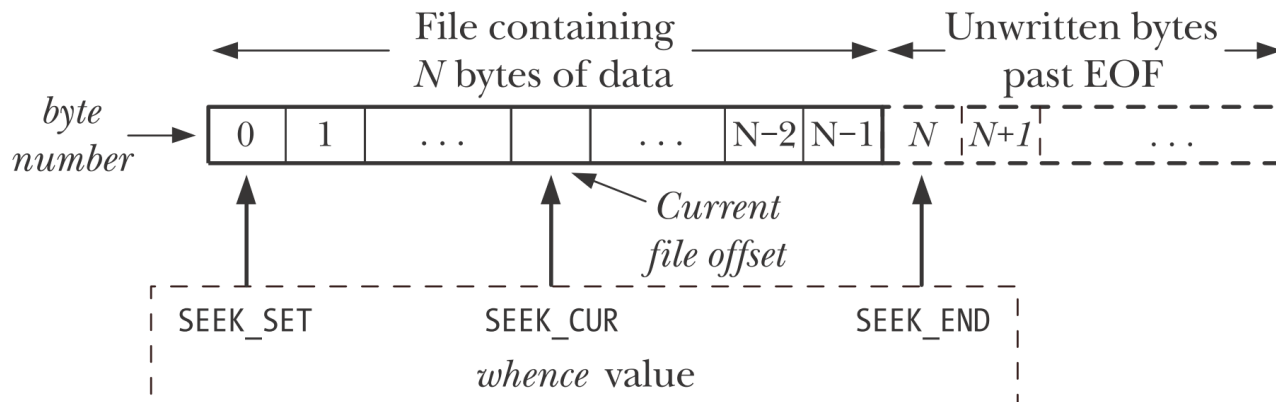
```
2  #include <unistd.h>
3  off_t lseek(int fd , off_t offset , int whence );
4  /*Returns new file offset if successful, or -1 on error*/
```

- It adjusts the file offset of the open file referred to by the *fd*, according to the values specified in *offset* and *whence*.
- Kernel records a *file offset* for each open file.
- This is the location in the file at which the next read or write will commence.
- When the file opened the offset is set to 0 i.e. the beginning of the file.

File I/O: lseek()



- *whence* argument can be
 - SEEK_SET
 - SEEK_CUR
 - SEEK_END



- `lseek()` simply adjusts the file offset, it doesn't cause any physical device access.
- We can't apply `lseek()` to pipe, FIFO, socket or terminal.

File I/O: lseek()



```
2 curr = lseek(fd, 0, SEEK_CUR); /* Retrives the file offset*/
3 lseek(fd, 0, SEEK_SET);        /* Start of file */
4 lseek(fd, 0, SEEK_END);        /* Next byte after the end of the file */
5 lseek(fd, -1, SEEK_END);       /* Last byte of file */
6 lseek(fd, -10, SEEK_CUR);      /* Ten bytes prior to current location */
7 lseek(fd, 10000, SEEK_END);    /* 10001 bytes past last byte of file */
```

- File holes
 - What if we read after `lseek(fd, 10000, SEEK_END)`?
 - Returns 0.
 - What if we write after `lseek(fd, 10000, SEEK_END)`?
 - It creates a file hole. File holes do not take disk space.
- File holes are useful when a program need to access a wide range of addresses (offset) but is unlikely to touch all of the potential blocks.
 - Virtual hard disks

Universality of I/O



- The same four system calls – *open()*, *read()*, *write()*, and *close()* – are used to perform I/O on all types of files.
 - Regular files, Pipe, FIFO, sockets, terminal devices
- *ioctl()* system call is for operations that fall outside the universal I/O model.

```
2  #include <sys/ioctl.h>
3  int ioctl(int fd , int request , ... /* argp */);
4  /*Value returned on success depends on request, or -1 on error*/
```

- *fd* refers to any file or device.
- *request* refers to the constant specific to the device.
- *argp* is the value or buffer depending the type of request.

Universality of I/O: *ioctl()*

- e.g. for updating flags of inode, *ioctl()* is used.

```
2 int attr;
3 if (ioctl(fd, FS_IOC_GETFLAGS, &attr) == -1)    /* Fetch current flags */
4     errExit("ioctl");
5 attr |= FS_NOATIME_FL; /*do not update last file access time*/
6 if (ioctl(fd, FS_IOC_SETFLAGS, &attr) == -1)    /* Update flags */
7     errExit("ioctl");
```

File Descriptors and Open Files



- It is possible to have multiple descriptors referring to the same open file.
- There are three data structures maintained by the kernel:
 - The per-process file descriptor
 - The system-wide table of open file descriptions
 - The file system i-node table.
- Per-process file descriptor table
 - Flags (close-on-exec)
 - A reference to the open file description

innovate **achieve** **lead**



File Descriptors and Open Files



- System-wide table of all open file descriptions
 - The current file offset
 - Modified by `read()`, `write()` and `lseek()`
 - Status flags (flags argument to `open()`)
 - File access mode (read-only, write only etc as specified in `open()`)
 - Settings related to signal driven I/O
 - a reference to i-node object for this file.
- i-node table
 - Each file system has a table of i-nodes for all files residing in the file system.
 - File type (regular file, socket, FIFO etc)
 - A pointer to list of blocks
 - Various properties of the file (size, timestamps etc)

File Descriptors and Open Files



- Two descriptors in different process may refer to the same open file entry.
 - `fork()`
 - Passing descriptor using UNIX domain sockets
- Two open file entries can refer to same i-node.
 - When the same is open twice in the same process or in different processes.
- When an open file entry is shared
 - Updating file offset or flags effects the other process.
- `close-on-exec` flag individual to a fd. Changing doesn't effect the other processes.

Duplicating File Descriptors: dup()



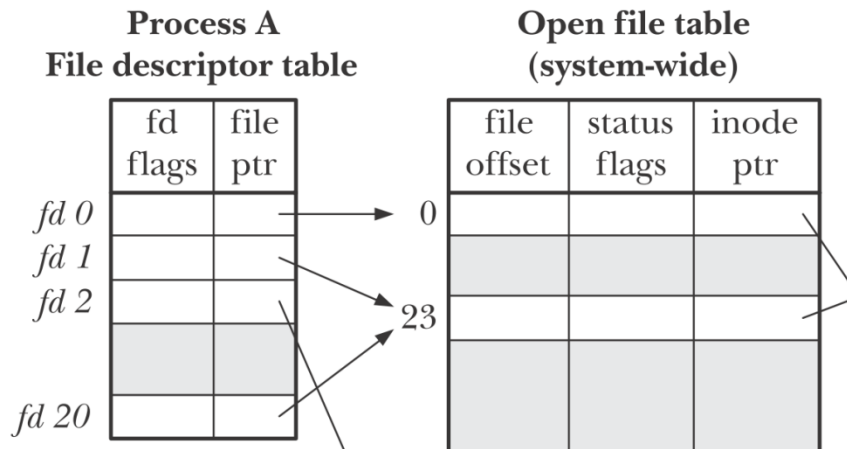
```
$ ./myscript > results.txt 2>&1
```

- Here 2>&1 indicates that standard error (fd 2) to be sent to the same place where standard output (fd 1) is being sent.
- This is possible by duplicating fd 2 to refer to open table entry referred by fd 1.

```
2  #include <unistd.h>
3  int dup(int oldfd );
4  /*Returns (new) file descriptor on success, or -1 on error*/
```

- Dup takes *oldfd* and returns a new fd that refers to the same open file table entry.
- New fd is guaranteed to be the lowest unused file descriptor.

Duplicating File Descriptors: dup2()



- `close(1); dup(20)`

```
2  #include <unistd.h>
3  int dup2(int oldfd , int newfd );
4  /*Returns (new) file descriptor on success, or -1 on error*/
```

- If *newfd* is open it closes it and copies the pointer in *oldfd* to *newfd* slot.
- This is done atomically.

File Control Operations: *fcntl()*



- `fcntl()` call performs operations on open file descriptors.

```
2  #include <fcntl.h>
3  int fcntl(int fd , int cmd , ...);
4  /*Return on success depends on cmd, or -1 on error*/
```

- `cmd` refers to commands.
- e.g. to change the flag after opening file

```
2  int flags;
3  flags = fcntl(fd, F_GETFL);
4  if (flags == -1)
5      errExit("fcntl");
6  flags |= O_APPEND;
7  if (fcntl(fd, F_SETFL, flags) == -1)
8      errExit("fcntl");
```

- Append is flag is being added to the flags in open file table entry.

File I/O Buffering



- Buffer cache:
 - Kernel reads the data from the disk and stores it in a buffer.
 - When a process request `read()`, the data is copied from kernel buffer to buffer in the user space.
 - Similarly when a user process writes, kernel writes to the buffer.
 - Kernel periodically syncs dirty buffers with disk.
- This allows `read()` and `write` to be faster.
- Use `fsync(int fd)` to forcefully sync buffers with disk.

Buffering in *stdio* Library

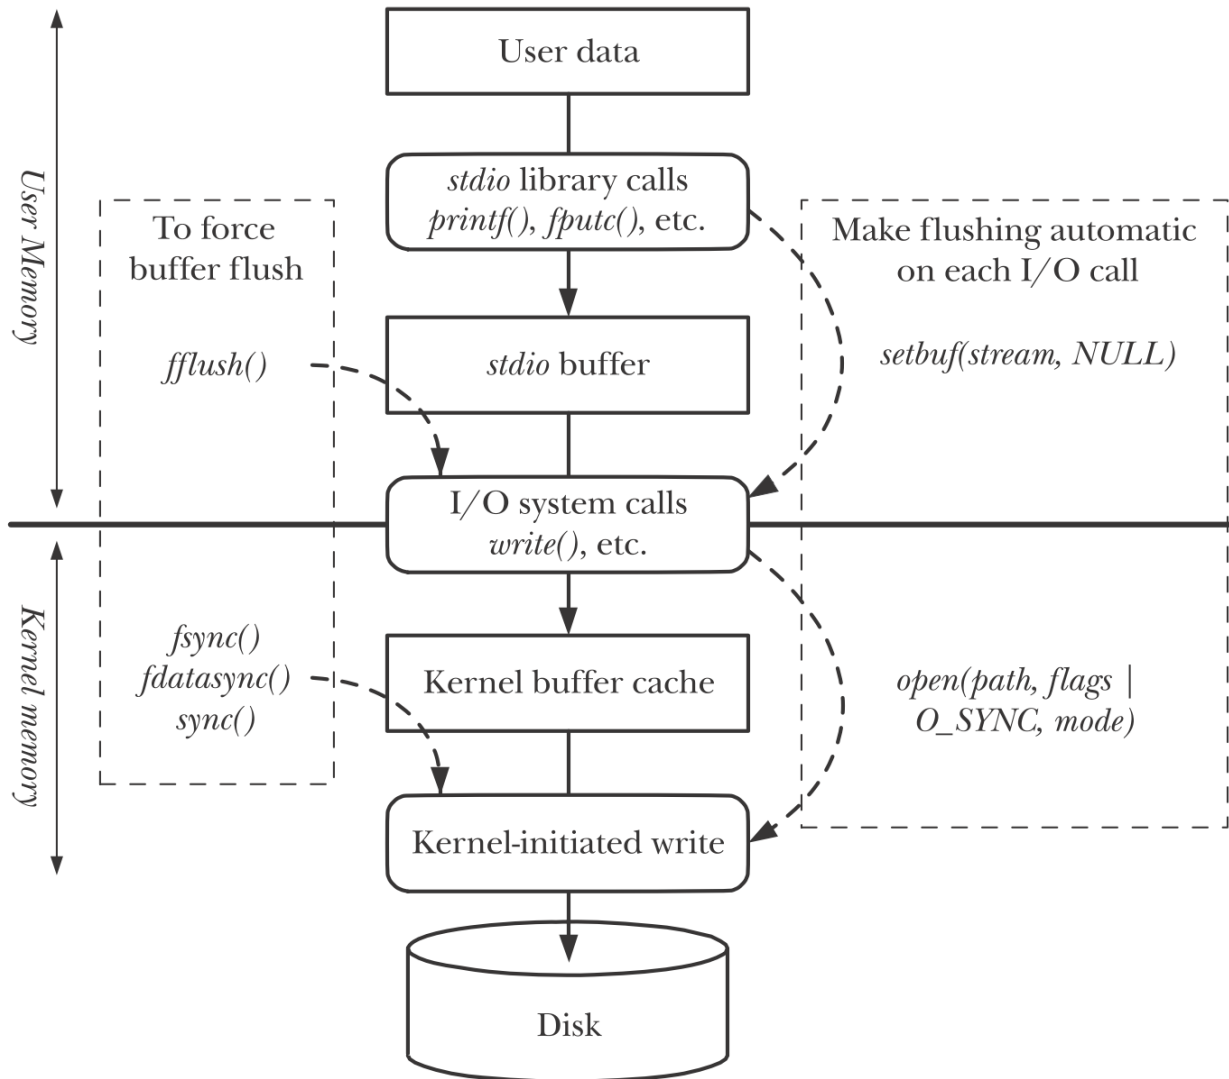


- C library buffers the data to reduce the number system calls (read, write).

```
2  #include <stdio.h>
3  int setvbuf(FILE * stream , char * buf , int mode , size_t size );
4  /*Returns 0 on success, or nonzero on error*/
```

- This is a library function that controls the type of buffering.
- This function must be called before any I/O operation.
- If *buf* is null, stdio automatically allocates the buffer for use with the *stream*.
- mode
 - `_IONBF`: no buffering. E.g. `stderr`
 - `_IOLBF`: line buffering. Default for terminal devices. Output is buffered until newline char. Data is read a line at a time.
 - `_IOFBF`: fully buffered I/O. data is read or written in units of buffer size. Default for disk files.
 - Use `fflush()` for flushing the buffer.

Buffering



Q&A



Next Time



- Please read through R1: chapters 6,7, 24-28



BITS Pilani
Pilani Campus



Thank You