



BITS Pilani
Pilani Campus

Network Programming

K Hari Babu
Department of Computer Science & Information Systems

Outline



- POSIX Message Queues
- POSIX Semaphores
- POSIX Shared Memory



Overview of POSIX IPC (R1: Ch51)

- Aim was to devise a set of IPC mechanisms that did not suffer the deficiencies of System V IPC.
 - Message Queues
 - Semaphores
 - Shared memory

API Overview



- All three mechanisms have similar API.

Table 51-1: Summary of programming interfaces for POSIX IPC objects

Interface	Message queues	Semaphores	Shared memory
Header file	<code><mqueue.h></code>	<code><semaphore.h></code>	<code><sys/mman.h></code>
Object handle	<code>mqd_t</code>	<code>sem_t *</code>	<code>int</code> (file descriptor)
Create/open	<code>mq_open()</code>	<code>sem_open()</code>	<code>shm_open() + mmap()</code>
Close	<code>mq_close()</code>	<code>sem_close()</code>	<code>munmap()</code>
Unlink	<code>mq_unlink()</code>	<code>sem_unlink()</code>	<code>shm_unlink()</code>
Perform IPC	<code>mq_send()</code> , <code>mq_receive()</code>	<code>sem_post()</code> , <code>sem_wait()</code> , <code>sem_getvalue()</code>	operate on locations in shared region
Miscellaneous operations	<code>mq_setattr()</code> —set attributes <code>mq_getattr()</code> —get attributes <code>mq_notify()</code> —request notification	<code>sem_init()</code> —initialize unnamed semaphore <code>sem_destroy()</code> —destroy unnamed semaphore	(none)

- Three stages
 - Create/open IPC object
 - Perform operations
 - Close/remove the object
- Create/open IPC object
 - Each IPC mechanism has an associated open call (mq_open(), sem_open(), or shm_open()),
 - which is analogous to the traditional UNIX open() system call used for files.

```
2 fd = shm_open("/mymem", O_CREAT | O_RDWR, S_IRUSR | S_IWUSR);
```

IPC Object Names



- Identifying a POSIX IPC object is via a name consisting of an initial slash, followed by one or more nonslash characters;
 - For example, /myobject.

Closing an IPC object



- *Close()* call that indicates that the calling process has finished using the object.
- IPC objects are automatically closed if the process terminates or performs an *exec()*.
- As with open files, POSIX IPC objects are reference counted—the kernel maintains a count of the number of open references to the object.
- Each IPC object has a corresponding *unlink* call whose operation is analogous to the traditional *unlink()* system call for files.
 - The *unlink* call immediately removes the object's name, and then destroys the object once all processes cease using it

Listing All IPC Objects



- Listed under the root directory (/), and the standard ls and rm commands can be used to list and remove IPC objects.
- Compiling programs that use POSIX IPC on Linux
 - On Linux, programs employing the POSIX IPC mechanisms must be linked with the realtime library, librt, by specifying the -lrt option to the cc command.

Comparison of System V IPC and POSIX IPC



- The POSIX IPC interface is simpler than the System V IPC interface.
- The POSIX IPC model—the use of names instead of keys, and the open, close, and unlink functions—is more consistent with the traditional UNIX file model.
- POSIX IPC objects are reference counted. This simplifies object deletion, because we can unlink a POSIX IPC object, knowing that it will be destroyed only when all processes have closed it.
- There is one notable advantage in favor of System V IPC: portability.

Message Queues



- POSIX message queues are reference counted.
- Each System V message has an integer type, and messages can be selected in a variety of ways using `msgrcv()`. By contrast, POSIX messages have an associated priority, and messages are always strictly queued (and thus received) in priority order.
- POSIX message queues provide a feature that allows a process to be asynchronously notified when a message is available on a queue.
- On Linux, POSIX message queues can be monitored using `poll()`, `select()`, and `epoll`. System V message queues don't provide this feature.
- POSIX message queues are less portable.
- The facility to select System V messages by type provides slightly greater flexibility than the strict priority ordering of POSIX messages.

Creating POSIX Message Queue



- The `mq_open()` function creates a new message queue or opens an existing queue.

```
2 #include <fcntl.h> /* Defines O_* constants */
3 #include <sys/stat.h> /* Defines mode constants */
4 #include <mqueue.h>
5 mqd_t mq_open(const char *name, int oflag, ...
6 /* mode_t mode, struct mq_attr *attr */);
7 //Returns a message queue descriptor on success, or (mqd_t) -1 on error
```

Table 52-1: Bit values for the `mq_open()` *oflag* argument

Flag	Description
O_CREAT	Create queue if it doesn't already exist
O_EXCL	With O_CREAT, create queue exclusively
O_RDONLY	Open for reading only
O_WRONLY	Open for writing only
O_RDWR	Open for reading and writing
O_NONBLOCK	Open in nonblocking mode

Closing a message queue



- The `mq_close()` function closes the message queue descriptor `mqdes`.

```
2  #include <mqueue.h>
3  int mq_close(mqd_t mqdes);
4  //Returns 0 on success, or -1 on error
```

- Removing a message queue

```
1  #include <mqueue.h>
2  int mq_unlink(const char *name);
3  //Returns 0 on success, or -1 on error
```

Differences with System V IPC



- Message Queues
- Semaphores
- Shared Memory

Q&A



Next Time



- Please read through R1: chapters 47-48



BITS Pilani
Pilani Campus



Thank You