


Preforked and prethreaded servers



- Traditional concurrent server model:
 - Fork a child after accepting a new client connection.
 - Good enough for low traffic services.
- For very high-load servers
 - web servers handling thousands of requests per minute
 - the cost of creating a new child (or even thread) for each client imposes a significant burden on the server.
- Instead of creating a new child process (or thread) for each client, the server precreates a fixed number of child processes (or threads) on startup.
 - Each child (thread) handles a new client. After completing one client, it accepts another connection.

- Different models
 - Child calls `accept()`
 - TCP Preforked Server, No Locking Around `accept`
 - TCP Preforked Server, Thread Locking Around `accept`
 - Parent calls `accept()` and passes the descriptor to child
 - TCP Preforked Server, Descriptor Passing

Preforking or Process Pool



Preforking or Process Pool

innovate

achieve

lead

```
static int      nchildren;
static pid_t    *pids;

int
main(int argc, char **argv)
{
    int          listenfd, i;
    socklen_t    addrlen;
    void          sig_int(int);
    pid_t        child_make(int, int, int);
    if (argc == 3)
        listenfd = Tcp_listen(NULL, argv[1], &addrlen);
    else if (argc == 4)
        listenfd = Tcp_listen(argv[1], argv[2], &addrlen);
    else
        err_quit("usage: serv02 [ <host> ] <port#> <#children>");
    nchildren = atoi(argv[argc-1]);
    pids = Calloc(nchildren, sizeof(pid_t));
    for (i = 0; i < nchildren; i++)
        pids[i] = child_make(i, listenfd, addrlen);    /* parent returns */

    Signal(SIGINT, sig_int);
    for ( ; ; )
        pause();    /* everything done by children */
}
```

Preforking or Process Pool: No Locking Around Accept

innovate

achieve

lead

```
pid_t
child_make(int i, int listenfd, int addrlen)
{
    pid_t    pid;
    void      child_main(int, int, int);
    if ( (pid = Fork()) > 0)
        return(pid);                /* parent */
    child_main(i, listenfd, addrlen); /* never returns */
}
/* end child_make */
/* include child_main */
void
child_main(int i, int listenfd, int addrlen)
{
    int                connfd;
    void               web_child(int);
    socklen_t          cliilen;
    struct sockaddr *cliaddr;
    cliaddr = Malloc(addrlen);
    printf("child %ld starting\n", (long) getpid());
    for ( ; ; ) {
        cliilen = addrlen;
        connfd = Accept(listenfd, cliaddr, &cliilen);
        web_child(connfd);           /* process the request */
        Close(connfd);
    }
}
/* end child main */
```



- **Advantages:**
 - No cost of fork() before responding to client.
 - Process control is simpler.
- **Disadvantages:**
 - Parent must guess how many children to fork.
 - If too less, clients will experience delays in response.
 - If too excessive, system performance degrades.